

Machine Learning Exercises: Set 1

Roger Garriga Calleja

January 30, 2017

Problem 5: Consider a binary classification problem in which the observation \mathbf{X} is real valued, $\mathbb{P}\{Y = 0\} = \mathbb{P}\{Y = 1\} = \frac{1}{2}$, and the class-oriented cumulative distribution functions are

$$\mathbb{P}\{X \leq x|Y = 0\} = \begin{cases} 0 & \text{if } x \leq 0 \\ \frac{x}{2} & \text{if } 0 < x \leq 2 \\ 1 & \text{if } x > 2 \end{cases} \quad \text{and} \quad \mathbb{P}\{X \leq x|Y = 1\} = \begin{cases} 0 & \text{if } x \leq 1 \\ \frac{(x-1)}{3} & \text{if } 1 < x \leq 4 \\ 1 & \text{if } x > 4 \end{cases} .$$

Determine $\eta(x) = \mathbb{P}\{Y = 1|X = x\}$. Compute the Bayes classifier and the Bayes risk R^* . Compute the asymptotic risk R_{1-NN} of the nearest neighbor classifier.

From the cdf given we can get the pdf just by taking derivatives, so

$$\mathbb{P}\{X = x|Y = 0\} = \begin{cases} 0 & \text{if } x \leq 0 \\ \frac{1}{2} & \text{if } 0 < x \leq 2 \\ 0 & \text{if } x > 2 \end{cases} \quad \text{and} \quad \mathbb{P}\{X = x|Y = 1\} = \begin{cases} 0 & \text{if } x \leq 1 \\ \frac{1}{3} & \text{if } 1 < x \leq 4 \\ 0 & \text{if } x > 4 \end{cases} .$$

So we can consider only $x \in [0, 4]$. Then, since Y is either 1 or 0, applying the law of total probabilities

$$\mathbb{P}(X = x) = \mathbb{P}(X = x|Y = 1)\mathbb{P}(Y = 1) + \mathbb{P}(X = x|Y = 0)\mathbb{P}(Y = 0) = \begin{cases} \frac{1}{4} & \text{if } 0 < x \leq 1 \\ \frac{5}{12} & \text{if } 1 < x \leq 2 \\ \frac{1}{6} & \text{if } 2 < x \leq 4 \end{cases}$$

. Using Bayes theorem

$$\eta(x) = \mathbb{P}\{Y = 1|X = x\} = \frac{\mathbb{P}(X = x|Y = 1)\mathbb{P}(Y = 1)}{P(X)} = \tag{1}$$

$$= \frac{\mathbb{P}(X = x|Y = 1)\mathbb{P}(Y = 1)}{\mathbb{P}(X = x|Y = 1)\mathbb{P}(Y = 1) + \mathbb{P}(X = x|Y = 0)\mathbb{P}(Y = 0)} . \tag{2}$$

Now, substituting on the equation

$$\eta(x) = \begin{cases} 0 & \text{if } x \leq 1 \\ \frac{2}{5} & \text{if } 1 < x < 2 \\ 1 & \text{if } 2 < x < 4 \end{cases} . \tag{3}$$

$$1 - \eta(x) = \begin{cases} 1 & \text{if } x \leq 1 \\ \frac{3}{5} & \text{if } 1 < x \leq 2 \\ 0 & \text{if } 2 < x \leq 4 \end{cases} . \tag{4}$$

The Bayes classifier takes the optimal decision as

$$p^*(X) = \begin{cases} 1 & \text{if } \eta(X) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases} \Leftrightarrow \begin{cases} 1 & \text{if } 2 < x \leq 4 \\ 0 & \text{otherwise} \end{cases} . \tag{5}$$

Now, we compute the Bayes risk as $R^*(x) = \mathbb{E}[\min\{\eta(x), 1 - \eta(x)\}]$. The minimum is

$$\min\{\eta(x), 1 - \eta(x)\} = \begin{cases} 0 & \text{if } x \leq 0 \\ \frac{2}{5} & \text{if } 1 < x \leq 2 \\ 0 & \text{if } 2 < x \leq 4 \end{cases} . \quad (6)$$

Finally, integrating $\min\{\eta(x), 1 - \eta(x)\}\mathbb{P}(X)$ over the space $[0, 4]$ to compute the expected value we get $R^* = \frac{1}{6}$.

For the 1-NN, the asymptotic risk is computed as $R_{1-NN} = 2\mathbb{E}[\eta(x)(1 - \eta(x))]$. The product is

$$\eta(x)(1 - \eta(x)) = \begin{cases} 0 & \text{if } x \leq 1 \\ \frac{6}{25} & \text{if } 1 < x \leq 2 \\ 0 & \text{if } 2 < x \leq 4 \end{cases} . \quad (7)$$

Then, integrating $\eta(x)(1 - \eta(x))\mathbb{P}(X)$ over the space $[0, 4]$ to compute the expected value we get $R_{1-NN} = \frac{1}{5}$.

Problem 6: Consider a binary classification problem in which both class-conditional densities are multivariate normal of the form

$$f_i(x) = \frac{1}{\sqrt{(2\pi)^d \det \Sigma_i}} e^{-\frac{1}{2}(x-m_i)^T \Sigma_i^{-1}(x-m_i)}, \quad i = 0, 1, \quad (8)$$

where $m_i = \mathbb{E}[X|Y = i]$ and Σ_i is the covariance matrix for class i . Let $q_0 = \mathbb{P}\{Y = 0\}$ and $q_1 = \mathbb{P}\{Y = 1\}$ be the a priori probabilities.

Determine the Bayes classifier. Characterize the cases when the Bayes decision is linear (i.e, it is obtained by thresholding a linear function of x).

To determine the Bayes classifier first we need the posterior probabilities,

$$\eta(x) = \mathbb{P}(Y = 1|X = x) = \frac{f_1(x)q_1}{f_1(x)q_1 + f_0(x)q_0}, \quad (9)$$

$$1 - \eta(x) = \mathbb{P}(Y = 0|X = x) = \frac{f_0(x)q_0}{f_1(x)q_1 + f_0(x)q_0}. \quad (10)$$

Then, $p^*(x) = \begin{cases} 1 & \text{if } \eta(x) \geq 1 - \eta(x) \\ 0 & \text{otherwise} \end{cases} = \begin{cases} 1 & \text{if } f_1(x)q_1 \geq f_0(x)q_0 \\ 0 & \text{otherwise} \end{cases}$. So we need to compare $f_1(x)q_1$ with $f_0(x)q_0$ and find the values of x for which $\eta(x) \geq 1 - \eta(x)$.

$$f_1(x)q_1 \geq f_0(x)q_0 \Leftrightarrow \frac{q_1}{\sqrt{\det \Sigma_1}} e^{-\frac{1}{2}(x-m_1)^T \Sigma_1^{-1}(x-m_1)} \geq \frac{q_0}{\sqrt{\det \Sigma_0}} e^{-\frac{1}{2}(x-m_0)^T \Sigma_0^{-1}(x-m_0)} \Leftrightarrow \quad (11)$$

$$\Leftrightarrow \log(q_1) - \frac{1}{2} \log(\det \Sigma_1) - \frac{1}{2}(x-m_1)^T \Sigma_1^{-1}(x-m_1) \geq \log(q_0) - \frac{1}{2} \log(\det \Sigma_0) - \frac{1}{2}(x-m_0)^T \Sigma_0^{-1}(x-m_0). \quad (12)$$

In general, the above inequation is quadratic, so the bounds would not be linear. To have a linear bound we need the second order term to be equal in both sides of the inequation, so we can get rid of it. Developing the quadratic product we get

$$(x-m_1)^T \Sigma_1^{-1}(x-m_1) \geq (x-m_0)^T \Sigma_0^{-1}(x-m_0) \Leftrightarrow \quad (13)$$

$$x^T \Sigma_1^{-1} x - x^T \Sigma_1^{-1} m_1 - m_1^T \Sigma_1^{-1} x + m_1^T \Sigma_1^{-1} m_1 \geq x^T \Sigma_0^{-1} x - x^T \Sigma_0^{-1} m_0 - m_0^T \Sigma_0^{-1} x + m_0^T \Sigma_0^{-1} m_0. \quad (14)$$

The second order term will be equal in both sides iff $\Sigma_1 = \Sigma_0$, so in this particular case the decision boundary will be linear.

Problem 7: Let the joint distribution of (X, Y) be such that X is uniform on the interval $[0, 1]$, and for all $x \in [0, 1]$, $\eta(x) = x$. Determine the prior probabilities $\mathbb{P}\{Y = 0\}$, $\mathbb{P}\{Y = 1\}$ and the class-conditional densities $f(x|Y = 0)$ and $f(x|Y = 1)$. Calculate R^* , R_{1-NN} and R_{3-NN} (i.e., the Bayes risk and the asymptotic risk of the 1-, and 3-nearest neighbor rules).

The posterior probabilities are $\eta(x) = x$ and $1 - \eta(x) = 1 - x$ and the distribution of X is $p(x) = \mathbb{1}_{x \in [0, 1]}$. From them we can get the prior probabilities by applying the law of total probabilities

$$\mathbb{P}(Y = 1) = \int_0^1 \mathbb{P}(Y = 1|x)p(x)dx = \int_0^1 xdx = \frac{1}{2} \quad (15)$$

$$\mathbb{P}(Y = 0) = 1 - \mathbb{P}(Y = 1) = \frac{1}{2}. \quad (16)$$

The class conditional densities can be calculated applying Bayes theorem

$$f(x|Y = 1) = \frac{\eta(x)\mathbb{P}(x)}{\mathbb{P}(Y = 1)} = 2x \quad (17)$$

$$f(x|Y = 0) = \frac{\eta(x)\mathbb{P}(x)}{\mathbb{P}(Y = 0)} = 2(1 - x) \quad (18)$$

Now, let's compute the risks. To compute the risk first we have to look for the minimum between $\eta(x)$ and $1 - \eta(x)$.

$$\min(\eta(x), 1 - \eta(x)) = \min(x, 1 - x) = \begin{cases} x & \text{if } 0 \leq x \leq \frac{1}{2} \\ 1 - x & \text{if } \frac{1}{2} < x \leq 1 \end{cases}. \quad (19)$$

With that we can compute the risks:

$$R^* = \mathbb{E}[\min(\eta(x), 1 - \eta(x))] = \int_0^{\frac{1}{2}} xdx + \int_{\frac{1}{2}}^1 1 - xdx = \frac{1}{4}, \quad (20)$$

$$R_{1-NN} = 2\mathbb{E}[\eta(x)(1 - \eta(x))] = 2 \int_0^1 x(1 - x)dx = \frac{1}{3}, \quad (21)$$

$$R_{3-NN} = \mathbb{E}[\eta(x)(1 - \eta(x))] + 4\mathbb{E}[\eta(x)^2(1 - \eta(x))^2] = \frac{3}{10}. \quad (22)$$

Problem 8 Write a program that generates training data of n i.i.d. pairs $(X_1, Y_1), \dots, (X_n, Y_n)$ of random variables distributed such that X takes values in \mathbb{R}^d and $Y \in \{0, 1\}$. The joint distribution is such that X is uniformly distributed in $[0, 1]^d$ and $\mathbb{P}\{Y = 1|X = x\} = x^{(1)}$ (where $x^{(1)}$ is the first component of $x = (x^{(1)}, \dots, x^{(d)})$).

Classify X using the 1, 3, 5, 7, 9-nearest neighbor rules. Re-draw (X, Y) many times so that you can estimate the risk of these rules. Try this for various values of n and d and plot the estimated risk. Explain what you observe.

The risk increases with the dimension d of x tending to 0.5. It was expected, since the response only depends on the first component of the vector $X = (X^{(1)}, \dots, X^{(d)})$, meaning that the rest of the components will add noise on the predictions. As the dimensionality increases, the relevance of a single component on the distance between two points is diluted. In general, the risk decreases with the number of neighbors used and the size of the data have little effect.




Figure 1: Risk of K-NN for $K=1,3,5,7,9$ with $N=50$

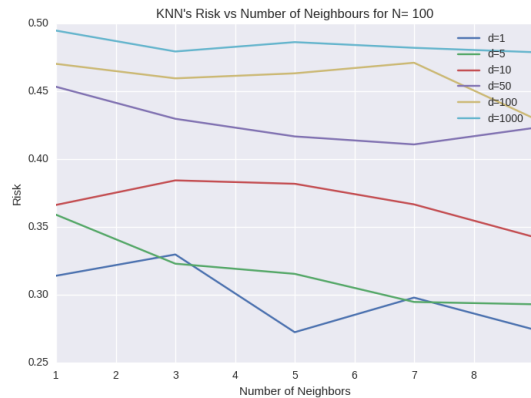


Figure 2: Risk of K-NN for $K=1,3,5,7,9$ with $N=100$




Figure 3: Risk of K-NN for $K=1,3,5,7,9$ with $N=1000$




Figure 4: Risk of K-NN for $K=1,3,5,7,9$ with $N=10000$

1 Appendix

Ex 8:

```
import sklearn as sk
import numpy as np
from numpy import random as rand
from sklearn import neighbors
import matplotlib.pyplot as plt
import seaborn as sns

def risk_knn(n,d,k,m):

    train_X=rand.uniform(low=0,high=1,size=(n,d)) #Generate X of uniform distribution
    train_p=train_X[:,0] #Generates Y from X as a Bernoulli with p=x[0] in each case
    train_Y=rand.binomial([1]*len(train_p),train_p,size=len(train_p))

    knn=neighbors.KNeighborsClassifier(n_neighbors=k) #Generate the KNN function
    knn.fit(train_X,train_Y) #Train the models

    test_X=rand.uniform(low=0,high=1,size=(m,d)) #Generate test data
    test_p=test_X[:,0]
    test_Y=rand.binomial([1]*len(test_p),test_p,size=len(test_p))

    knn_pred=knn.predict(test_X) #Predict values

    knn_check=test_Y==knn_pred #Compute error frequency
    knn_wrong=m-sum(knn_check)
    return knn_wrong/m

n=[10000]*5
m=[10000]*5
k=[1,3,5,7,9]

#Estimate risk knn for different d
i=1
all_risks=[0]*6
for d in [1,5,10,50,100,1000]:

    s=[d]*5
    risks=list(map(risk_knn,n,s,k,m))
    all_risks[i-1]=risks
    #plt.subplot(3,2,i)
    plt.plot(k,risks)
    #plt.ylim([0.15,0.5])
    plt.title("KNN's Risk vs. 'K'NN. N=%0.0f D=%0.0f" %(n[0],d))
    i+=1
```

```
plt.legend(['d=1', 'd=5', 'd=10', 'd=50', 'd=100', 'd=1000'], loc='upper right')
plt.xlabel('Number of Neighbors')
plt.ylabel('Risk')
plt.title("KNN's Risk vs Number of Neighbours for N=%0f" %n[0])
```